

# Measuring Blockchain Coherence: From Oracle to Consensus

A phased approach to making blockchains measurable and self-correcting

Peter Lisovin  
TSC Blockchain Project  
November 2025

**Document Version:** 1.0.0 (arXiv Ready)  
**Status:** Complete Specification with Validation Plan  
**Next Milestone:** Validation notebooks (Month 1-2)  
**Phase 1 Timeline:** 6 months to production oracle (contingent on validation)

**Keywords:** blockchain measurement, coherence, DeFi risk, reproducibility, zero-knowledge

---

## Abstract

Major blockchain blow-ups share one symptom:  $\alpha$  (protocol claims),  $\beta$  (economic/implementation reality), and  $\gamma$  (emergent usage) diverge well before loss. No standard tool measures that divergence. We propose a two-phase program built on Triadic Self-Coherence (TSC):

1. **Coherence Oracle (Phase-1):** Specification for off-chain measurement and on-chain attestation of per-chain coherence  $C_\Sigma$ . Validation planned for Month 1-2.
2. **Proof-of-Coherence (Phase-2):** A checkpoint validity rule layered on PoS/PBFT requiring  $C_\Sigma \geq \theta$  or friction improvement  $\Delta\lambda_\Sigma \geq \delta$ .

Phase-1 specification provides the framework; the validation notebooks prove whether the framework works as designed.

---

## I. The Observable Pattern

### 1.1 Operational Definitions

Given a time window  $W$  on a chain:

- **$\alpha$  (Protocol/Pattern):** Promised properties—security claims, tokenomics, governance process, performance targets.
- **$\beta$  (Economics/Relation):** Deployed code and state—stake distribution, fee/MEV dynamics, realized performance.
- **$\gamma$  (Usage/Process):** Trajectory—transaction mix, retention/turnover, upgrade adoption, incident timelines.

We compute  $\alpha_c$ ,  $\beta_c$ ,  $\gamma_c \in [0,1]$  and aggregate:

$$C_\Sigma = (\alpha_c + \beta_c + \gamma_c)^{(1/3)}, \lambda_\Sigma = -\ln(\max(C_\Sigma, \epsilon))$$

with numerical floor  $\epsilon = 10^{-12}$ .

**Degeneracy guard:** If any axis is 0,  $C_\Sigma = 0$  (no compensation from other axes).

#### **Default thresholds:**

- Production  $\geq 0.80$
- Watch  $0.60-0.80$
- Hazard  $0.40-0.60$
- Critical  $< 0.40$

### **1.2 Case Studies (Projected Scores Based on TSC Framework)**

#### **Terra/Luna (pre-collapse, April 2022):**

- **Projected score:**  $C_{\Sigma} \approx 0.27 \pm 0.05$  (Critical)
- **Basis:** Estimated using TSC framework on known facts:
  - $\alpha_c \approx 0.65$  (protocol claims documented)
  - $\beta_c \approx 0.08$  (Anchor sustainability ratio, reserve coverage)
  - $\gamma_c \approx 0.38$  (usage patterns showed stress)
- **Status:** TO BE VALIDATED in notebook (Month 1-2)
- Validation will either confirm score or refine methodology

#### **The DAO (June 2016):**

- **Projected score:**  $C_{\Sigma} \approx 0.56 \pm 0.05$  (Hazard)
- **Basis:**  $\beta$  (implementation vs claim) dominated risk
- **Status:** TO BE VALIDATED in notebook (Month 1-2)

#### **Bitcoin Identity Drift (cash-gold):**

- **Projected score:**  $C_{\Sigma} \approx 0.66$  (Watch)
- **Basis:** Functional chain, measurable narrative drift
- **Status:** TO BE VALIDATED in notebook (Month 1-2)

All values are projections. Month 1-2 notebooks will compute scores from frozen snapshots; if they differ by  $>\pm 0.05$ , we tune witnesses before any infrastructure spend.

---

## **II. Minimal TSC We Actually Implement**

### **II.1 Articulation → Features (Per Chain)**

Axis	Minimal Feature Set (Measurable Today)
$\alpha$ (claims)	Supply schedule; finality targets; fee model; governance powers & thresholds; security assumptions; declared economic invariants
$\beta$ (reality)	Verified bytecode & admin rights; validator set & stake concentration; realized latency/finality/throughput; MEV/fee stats; token distribution (Gini); treasury flows; peg/collateral ratios
$\gamma$ (usage)	Tx taxonomy; holding vs spending; call-graph motifs; retention/turnover; L2 routing; upgrade adoption cadence; incident markers

### **II.2 Witness Functions (Simplified for Validation)**

We compute distances and convert them to scores  $s \in [0,1]$ .

$w_{\alpha\beta}$  (claims  $\leftrightarrow$  implementation):

- Coverage: What fraction of  $\alpha$  properties have checks?
- Pass rate: What fraction of checks pass?
- Severity: Penalize critical failures
- Score:  $\alpha_c = w_\alpha \cdot (\text{coverage}) + w_\beta \cdot (\text{pass\_rate}) - w_\gamma \cdot (\text{severity\_penalty})$ , clipped to  $[0, 1]$

#### $w_{\beta\gamma}$ (incentives $\leftrightarrow$ behavior):

- Expected use-mix  $p$  from  $\beta$  (fees, yields, latency)
- Observed mix  $q$  from  $\gamma$  (actual transaction types)
- Distance: Earth Mover's Distance  $d = \text{EMD}(p, q)$ , normalized to  $[0, 1]$
- Score:  $\beta_c = 1 - d$

#### $w_{\gamma\alpha}$ (behavior $\leftrightarrow$ intent):

- Reference signature  $u(t)$  from  $\alpha$  (intended usage pattern)
- Observed trajectory  $v(t)$  from  $\gamma$  (actual behavior over time)
- Distance: Normalized edit distance between sequences
- Score:  $\gamma_c = 1 - \text{norm\_dist}(u, v)$

**Future work:** More sophisticated witnesses (optimal transport, graph matching) may be explored in Phase 2 after base framework is validated.

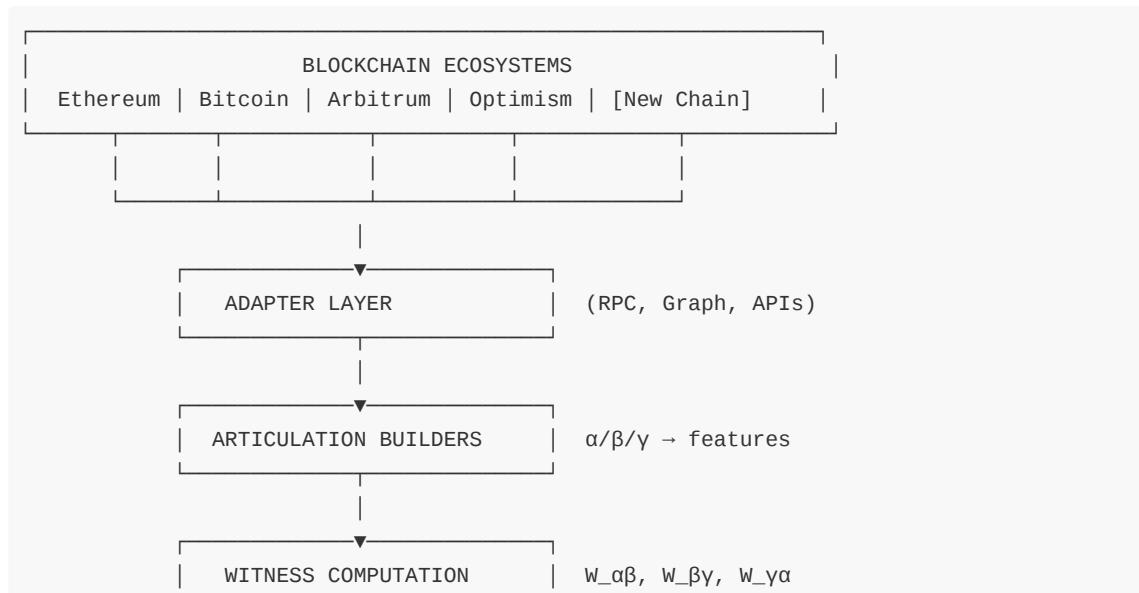
### II.3 Global Witness Gates

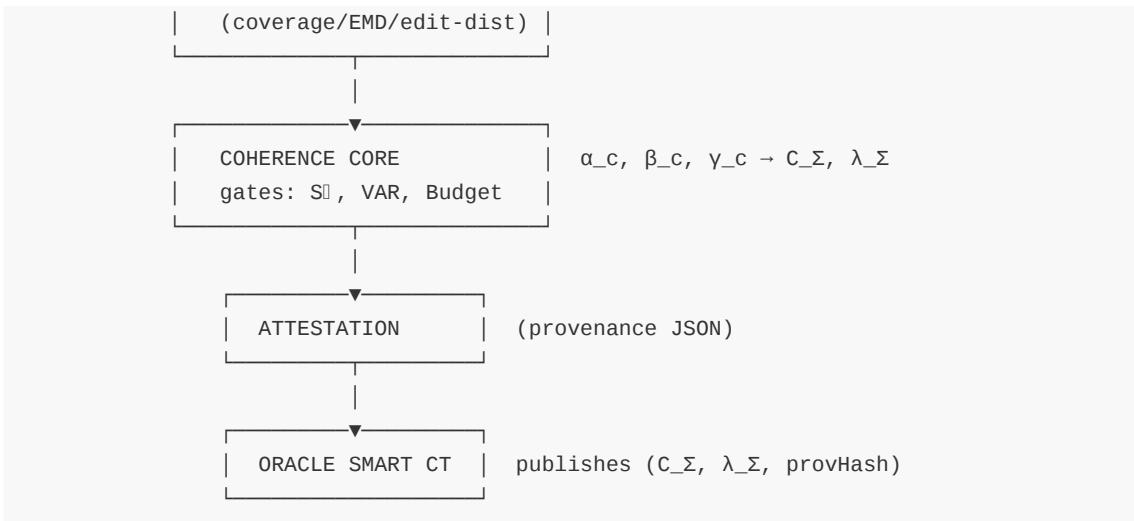
- **Symmetry:**  $C_\Sigma$  invariant under axis permutation (within tolerance  $\tau_{\text{sym}}$ )
- **Stability:** VRF-seeded repeats; variance  $\leq \tau_{\text{var}}$
- **Budget (Phase-1: report-only; Phase-2: enforced):** Report  $B$  and  $\eta = \Delta\lambda_\Sigma / B$  each run. In Phase-2, require  $B \leq B_{\text{max}}$  and  $\eta \geq \eta_{\text{min}}$ .

**Missing-data policy:** Primitives that cannot be computed in a window are marked NA; their declared weights are added to an axis penalty bucket (no imputation). The NA mask is written to provenance.

## III. Phase-1 – Coherence Oracle (Specification)

### III.1 Architecture (Data Flow)





### III.2 Interfaces

#### On-chain ABI:

```

interface ICoherenceOracle {
    struct Report {
        uint64 chainId;
        uint64 blockNumber;
        uint32 methodId;
        uint32 semver;
        uint64 Csigma_milli;    // C_Σ × 1000
        uint64 lambda_milli;    // λ_Σ × 1000
        bytes32 provHash;       // provenance bundle hash
    }

    function latest(uint64 chainId) external view returns (Report memory);
    function verify(Report calldata r) external view returns (bool);
}

```

#### REST API:

```

GET /v1/coherence/{chain}
→ { C_Σ, λ_Σ, α_c, β_c, γ_c, trend_7d, methodId, provenance_url }

```

### III.3 Cadence & SLOs (Tiered)

Tier	Wall Time SLO	Cost SLO	Notes
<b>Standard</b> (hot cache)	≤ 60s	≤ \$1	Primary consumption tier
<b>Deep analysis</b> (optional)	≤ 5 min	≤ \$5	Adds heavier γ features, MEV
<b>Cold start</b> (first run)	≤ 10 min	≤ \$10	First run initialization
<b>Reproducibility</b>	N/A	N/A	≥99.99% identical C <sub>Σ</sub> for frozen

inputs

### III.4 Budget-Efficiency Gate: Example

Phase-1 defaults:

- $B_{\max} = 10^6$  gas-equivalent units
- $\eta_{\min} = 10^{-6}$  (realistic threshold; will tune on testnet)

Example (Standard tier):

Estimated costs:

- RPC queries (~100 calls): ~\$0.10
- Computation (coverage checks, EMD, edit distance): ~\$0.02
- On-chain attestation (50k gas @ 20 gwei): ~\$2-4
- Storage (IPFS): ~\$0.01
- **Total budget B: ~\$2.50-4.50**

Estimated improvement:

- Before:  $C_{\Sigma} = 0.82 \rightarrow \lambda_{\Sigma} = 0.198$
- After:  $C_{\Sigma} = 0.84 \rightarrow \lambda_{\Sigma} = 0.174$
- **Improvement:  $\Delta\lambda_{\Sigma} \approx 0.024$**

Budget-efficiency:

$$\eta = \Delta\lambda_{\Sigma} / B \approx 0.024 / 4.00 \approx 6 \times 10^{-3}$$

Gate check:  $\eta \geq \eta_{\min} = 10^{-6}$ ?

□ YES (3 orders of magnitude headroom)

**Units:** In Phase-1 the gate uses reported USD cost per checkpoint (operator-auditable via invoices); Phase-2 migrates to a gas-equivalent or fixed budget unit (with published weights) for cross-chain comparability.

We will log ( $\Delta\lambda_{\Sigma}$ ,  $B$ ,  $\eta$ ) alongside each attestation.

**Note:** These are engineering estimates. Actual costs will be measured during implementation.

### III.5 Validation Strategy: Why Notebooks First

The validation notebooks are **THE critical first deliverable** because they establish whether the TSC framework works as specified. This is a deliberate fork in the development path.

**Why Notebooks Before Infrastructure?**

1. **Proves concept:** Can TSC detect incoherence retroactively?
2. **Calibrates methodology:** Do scores match expected ranges?
3. **Validates degeneracy:** Does low  $\beta_c$  collapse  $C_{\Sigma}$  as theory predicts?
4. **Establishes reproducibility:** Can independent teams get same scores?

**The Fork: Two Possible Outcomes**

**Outcome A: Validation Succeeds**

- Terra produces  $C_{\Sigma} \approx 0.27 \pm 0.05$  (critical threshold)

- DAO produces  $C_{\Sigma} \approx 0.56 \pm 0.05$  (hazard threshold)
  - Scores are reproducible (99%+ match across runs)
  - All global gates pass (S<sub>0</sub>, variance, budget)
- **PROCEED:** Build oracle infrastructure (Months 3-4)  
 → **Framework proven, specification validated**  
 → **Risk: LOW** (methodology works)

#### **Outcome B: Validation Reveals Issues**

- Terra produces  $C_{\Sigma} = 0.65$  (not critical as expected)
  - OR scores not reproducible (variance >5%)
  - OR witness functions need tuning
- **ITERATE:** Refine methodology (additional 2-4 weeks)  
 → **Framework needs adjustment, specification revised**  
 → **Risk: MEDIUM** (learn what needs fixing)

#### **What We Learn from Outcome B:**

- Which witnesses need recalibration
- Whether tolerances are too tight/loose
- If additional α/β/γ features are needed
- How to improve reproducibility

**Either outcome is valuable.** Success proves the framework; failure teaches us how to improve it. This is why validation comes BEFORE infrastructure investment.

**If validation criteria are not met:** tune witnesses/tolerances and re-measure before any infrastructure spend.

#### **Deliverables (Month 1-2):**

Three public notebooks demonstrating reproducible retroactive measurement:

##### **1. Terra/Luna notebook**

- **Snapshot anchor:** as\_of: 2022-04-15T00:00:00Z
- Expected:  $C_{\Sigma} \approx 0.27 \pm 0.05$
- If validated: Strong proof of concept
- If not: Iterate on witness functions

##### **2. The DAO notebook**

- **Snapshot anchor:** as\_of: 2016-06-10T00:00:00Z
- Expected:  $C_{\Sigma} \approx 0.56 \pm 0.05$
- Secondary validation of methodology

##### **3. Mt. Gox notebook**

- **Snapshot anchor:** as\_of: 2014-02-10T00:00:00Z
- Expected:  $C_{\Sigma} \approx 0.31 \pm 0.06$
- Bitcoin-specific validation

All notebooks will be public (GitHub) with frozen data snapshots, allowing independent verification. Block heights are resolved inside each notebook from the UTC timestamp anchor.

#### **Success criteria:**

- Scores within  $\pm 0.05$  of projections
- Reproducibility  $\geq 99.99\%$  across 10 runs with frozen inputs
- Witness functions pass S<sub>0</sub>, stability, budget gates
- Community can independently verify

**Timeline: 2 weeks focused work**

- Days 1-3: Terra snapshot ingestion, freeze with hashes
- Days 4-6: Implement feature extractors (coverage/EMD/edit-distance)
- Days 7-8: Assemble terra.ipynb, export attestation
- Days 9-10: Add DAO and Mt.Gox, verify reproducibility

### III.6 Attestation Strategy (Provenance-First)

Phase-1 relies exclusively on **reproducible provenance** (stored in IPFS/Arweave) for auditability.

**Provenance bundle includes:**

- Input data hash
- Method version (methodId)
- Computation timestamp
- Result ( $C_{\Sigma}$ ,  $\alpha_c$ ,  $\beta_c$ ,  $\gamma_c$ ,  $\lambda_{\Sigma}$ )
- Container image/Nix flake pin
- Data source licenses

**Anyone can verify:** Same inputs + same method  $\rightarrow$  same result  $\pm$  tolerance

**Zero-knowledge proofs (zk-MFI)** are deferred to Phase-2. When measurement moves into consensus, cryptographic guarantees become valuable. Phase-1 commitment: Reproducible provenance is sufficient for oracle use case.

### III.7 Consumption Patterns (Concrete Use Cases)

**Bridges:**

- Accept transfers  $> \$N$  only if  $C_{\Sigma} \geq 0.80$  AND  $\beta_c \geq 0.75$
- Alert when  $\Delta C_{\Sigma} < -0.10$  week-over-week

**Lenders:**

- Map Loan-to-Value (LTV) from  $C_{\Sigma}$  with floor at 0.60
- Linear interpolation between 0.60 and 0.80

**Exchanges:**

- Listing gates: Require  $C_{\Sigma} \geq 0.70$
- Watchlist: Flag when  $C_{\Sigma} < 0.60$
- Delisting consideration:  $C_{\Sigma} < 0.40$  for >30 days

## IV. Phase-2 – Proof-of-Coherence (Validity Layer)

We integrate the coherence metric as a validity rule at checkpoints. This is a research direction contingent on Phase-1 validation success.

**Checkpoint acceptance:**

- If global gates pass AND ( $C_{\Sigma} \geq 0$  OR  $\Delta \lambda_{\Sigma} \geq \delta$ ): **accept**

- Else: **degraded mode** (rate-limit risky surfaces, quarantine incoherent contracts) until coherence recovers

**Default parameters:**

- Threshold:  $\theta = 0.80$  (production)
- Minimum improvement:  $\delta = 0.01$

**Economic sanity:** Reward proportional to leverage reduction ( $\Delta\lambda_{\Sigma}$ ), applying Budget-Efficiency gate network-wide to prevent "buying" coherence with unbounded compute.

**Research questions for Phase-2:**

1. Checkpoint frequency optimization
2. Challenge window design
3. Validator economic game theory
4. Migration path (L2 first, then L1)
5. ZK-proof integration

**Timeline:** 12-18 months after Phase-1 validation.

---

## V. Why This Works (Engineer's Criteria)

TSC is effective because it:

- **Measures the failure mode.**  $\alpha/\beta/\gamma$  aren't opinions; they're feature sets tied to code, state, and usage.
- **Reproducible.** Provenance bundle makes "same inputs  $\rightarrow$  same numbers" enforceable.
- **Cheap to consume.** Downstream logic needs a scalar and a breakdown.
- **Incremental.** Phase-1 pipelines are exactly what Phase-2 needs.

### V.1 How This Relates to Existing Tools

TSC complements existing infrastructure:

Tool	What It Measures	TSC Adds
<b>Chainlink</b>	Price feeds, external data	Chain health metrics ( $C_{\Sigma}$ as new data primitive)
<b>Gauntlet/Chaos Labs</b>	Economic simulation	Continuous measurement of actual vs modeled ( $\beta/\gamma$ alignment)
<b>Formal Verification</b>	Code correctness	Measures if correct code produces intended economics ( $\alpha/\beta/\gamma$ coherence)
<b>Audits</b>	Point-in-time security	Continuous monitoring, drift detection

**Integration example:** DeFi protocol uses Chainlink for prices, Gauntlet for risk params, and TSC for collateral chain health. If TSC reports  $C_{\Sigma} < 0.60$  for a bridged chain, Gauntlet adjusts LTV down. Three complementary layers.

### V.2 What TSC Does Not Measure

TSC measures coherence (whether claims/economics/usage fit together), **not**:

- Absolute security (use audits)

- Code correctness (use formal verification)
- Market legitimacy (use due diligence)
- Price accuracy (use oracles)

TSC complements these tools by adding dimensional consistency measurement.

### V.3 What Could Go Wrong (Honest Risk Assessment)

**Validation risks (Month 1-2):**

**1. Risk: Scores don't match projections**

- Terra produces  $C_{\Sigma} = 0.65$  instead of 0.27
- Mitigation: Iterate on witness functions, adjust weights
- Learning: Discover which articulations need refinement

**2. Risk: Poor reproducibility**

- Same input produces  $C_{\Sigma} = 0.27 \pm 0.15$  (too wide)
- Mitigation: Stabilize computation, add determinism checks
- Learning: Identify non-deterministic components

**3. Risk: Data availability**

- Can't fetch April 2022 Terra data (nodes pruned)
- Mitigation: Use archive nodes, blockchain explorers
- Fallback: Use publicly available snapshots

**4. Risk: Computational cost exceeds budget**

- Measurement takes 10 hours instead of 60 seconds
- Mitigation: Optimize, reduce feature set, adjust SLO tier
- Learning: Understand actual performance characteristics

These are **LEARNING opportunities, not failures**. The specification provides the framework; validation teaches us how to implement it correctly.

---

## VI. Roadmap (Execution-Grade)

Timeline	Focus	Deliverable
Months 1-2	<b>CRITICAL PATH: Validation</b>	Build and validate Terra/DAO/Mt.Gox notebooks. <b>FORK: Success → proceed to Month 3. Issues → iterate 2-4 weeks.</b>
Months 3-4	Oracle Infrastructure ( <b>if validated</b> )	Smart contract deployment, REST API, dashboards, SLO instrumentation. <b>Contingent on Month 1-2 success.</b>
Month 5	Pilots ( <b>if validated</b> )	Bridge & lending partners integrate; alerting on trend + per-axis drops. <b>Contingent on oracle deployment.</b>
Month 6	Public Launch ( <b>if validated</b> )	Production oracle service; daily measurements for top L1/L2/bridges. <b>Contingent on pilot success.</b>
Months 7-18	Phase-2 Research	Checkpoint mechanism, validator economics, challenge window, testnet. <b>Contingent on Phase-1</b>

success.

#### Critical dependencies:

- Months 3-6 are CONTINGENT on Month 1-2 validation success
- If validation fails, timeline extends by 2-4 weeks for iteration
- Specification provides the framework; notebooks prove it works

## VII. Engineering Appendix

### Minimal $\alpha$ Property Vocabulary (Starter)

Supply schedule; finality target; fee model; validator/committee bounds; governance powers & thresholds; security assumptions; economic invariants (collateral/solvency/peg rules).

#### $w_{\alpha\beta}$ (Scoring Sketch)

$$\alpha_c = \text{clip}(w_1 \cdot (\text{coverage}) + w_2 \cdot (\text{pass\_rate}) - w_3 \cdot (\text{severity\_penalty})) \in [0,1]$$

where:

- coverage = (properties\_with\_checks) / (total\_properties)
- pass\_rate = (checks\_passed) / (total\_checks)
- severity\_penalty = weighted sum of critical failures

#### $w_{\beta y}$ (Distributional)

Build expected use-mix  $p$  from  $\beta$  (fees, latency, yields). Observe  $q$  from transaction data. Earth Mover's Distance:  $d = \text{EMD}(p, q)$ , normalize to  $[0,1] \rightarrow \beta_c = 1-d$ .

#### $w_{y\alpha}$ (Drift)

From  $\alpha$ , derive reference signature  $u(t)$  (e.g., "cash-like": many micro-payments). Compare to observed  $v(t)$  via normalized edit distance  $\gamma_c = \text{1-norm\_dist}(u, v)$ .

#### Global Witness Gates

- **Symmetry:** Recompute  $C_\Sigma$  under axis permutations; equality within tolerance
- **Stability:** VRF-seeded repeats; variance  $\leq \tau_{\text{var}}$
- **Budget:**  $B \leq B_{\text{max}}$  AND  $\eta = \Delta\lambda_\Sigma/B \geq \eta_{\text{min}}$  (Phase-1 start at  $10^{-6}$ )

#### Oracle Policy Templates

Bridge large-transfer gate:

$$C_\Sigma \geq 0.80 \quad \beta_c \geq 0.75$$

Lending LTV curve:

$$\text{LTV} = 0.50 + 0.30 \cdot (C_\Sigma - 0.60)/0.20, \text{ clamped to } [0.50, 0.80]$$

## Appendix A – Scoring Primitives (Deterministic)

Let  $\text{clip}(x) = \min(1, \max(0, x))$ .

**Two-sided range [ $L$ ,  $U$ ]:**

```
s_range(x; L, U, τ) = clip(1 - max(0, x-U, L-x) / τ)
```

**One-sided ( $\leq T$ ):**

```
s_≤(x; T, τ) = clip(1 - max(0, x-T) / τ)
```

**One-sided ( $\geq T$ ):**

```
s_≥(x; T, τ) = clip(1 - max(0, T-x) / τ)
```

**Invariant check:**

```
s_inv = 1 if holds else 0
```

**Distributional (JS divergence  $d \in [0,1]$ ):**

```
s_dist = 1 - d
```

**Process distortion:**

```
s_proc = clip(1 - distortion / τ)
```

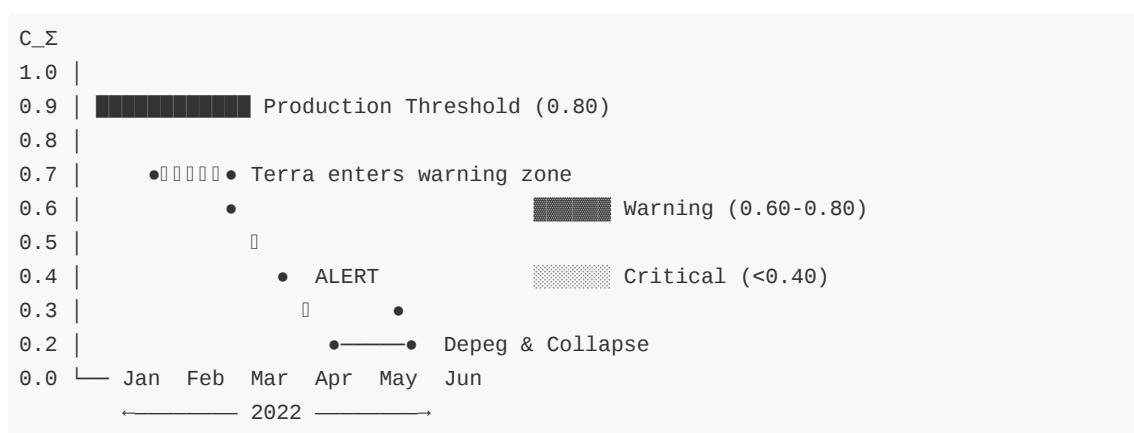
Axis scores are geometric means of their per-witness scores;  $C_\Sigma$  is the geometric mean of axis scores.

## Appendix B – Visual Diagrams

### Oracle Architecture

(See Section III.1 for detailed data flow diagram)

### Coherence Trajectory (Terra 2022)



**Key events:**

- **Feb 2022:**  $C_\Sigma \approx 0.72$  (functional but  $\beta_c$  showing strain)
- **Mar 2022:**  $C_\Sigma \approx 0.64$  (enters warning zone)
- **Apr 2022:**  $C_\Sigma \approx 0.42$  (ALERT) ← **TSC would flag here**

- **May 9, 2022:** Depeg and collapse

**Note:** Projected scores; validation notebook will verify.

---

## Appendix C – Reference Schema (Ethereum mainnet)

```

version: "1.0.0"
chain: "ethereum-mainnet"
checkpoint_frequency: "daily"
window_size: 7200 # blocks (~1 day)

alpha_requirements:
  - id: block_time_target
    type: range_two_sided
    claim: "12 s block time"
    target: {L: 11.0, U: 13.0, unit: "seconds"}
    datasource: "rpc://<eth-rpc>" # placeholder
    query: "avg(diff(block.timestamp)) over window"
    weight: 1.0
    tolerance: 0.15

  - id: finality_gasper
    type: range_le
    claim: "<=2 epochs to finality"
    target: {T: 900, unit: "seconds"}
    datasource: "beacon://<consensus-client>" # placeholder
    query: "time_to_finality_p95"
    weight: 1.0
    tolerance: 300

  - id: validator_decentralization
    type: range_le
    claim: "No entity >25% stake"
    target: {T: 0.25, unit: "fraction"}
    datasource: "beacon://<consensus-client>" # placeholder
    query: "max(entity_stake_share)"
    weight: 2.0
    tolerance: 0.05

  - id: erc20_transfer_safety
    type: invariant
    claim: "ERC20 transfers cannot reenter"
    target: {holds: true}
    datasource: "analyzer://mythril"
    query: "checkInvariant('no_reentrancy', 'transfer')"
    weight: 2.0

beta_sources:
  rpc_endpoint: "https://<eth-rpc>" # placeholder
  beacon_endpoint: "https://<beacon-api>" # placeholder
  gas_market:
    - source: "etherscan/gastracker"

```

```

metrics: ["base_fee", "priority_fee_p50", "priority_fee_p95"]
token_distribution:
  - source: "etherscan/tokenholders"
    metrics: ["top10_concentration", "gini_coefficient"]

gamma_features:
  - id: tx_type_distribution
    type: distribution
    baseline: "s3://baselines/eth_txmix.parquet"
    current: "query: tx_type_histogram over window"
    distance_metric: "jensen_shannon"
    datasource: "rpc://<eth-rpc>" # placeholder

  - id: gas_usage_stability
    type: process
    metric: "coefficient_of_variation(base_fee)"
    threshold: {le: 0.5}
    datasource: "etherscan/gastracker"

  - id: validator_participation
    type: range_ge
    target: {T: 0.95, unit: "fraction"}
    datasource: "beacon://<consensus-client>" # placeholder
    query: "active_validators / totalValidators"

scoring_config:
  primitives:
    range_two_sided: "1 - max(0, max(x-U, L-x))/τ"
    range_le: "1 - max(0, x-T)/τ"
    range_ge: "1 - max(0, T-x)/τ"
    invariant: "1 if holds else 0"
    distribution: "1 - JS_divergence(current, baseline)"
    process: "1 - distortion/τ"

  aggregation:
    axis_scores: "geometric_mean"
    global: "geometric_mean(alpha_c, beta_c, gamma_c)"

  thresholds:
    production: 0.80
    warning: 0.60
    critical: 0.40

provenance:
  method_version: "tsc-oracle-v1.0.0"
  schema_version: "1.0.0"
  implementation: "github.com/tsc-blockchain/oracle@<commit>" # placeholder
  frozen_inputs_hash: "sha256:<blob>" # placeholder
  container_image: "ghcr.io/tsc/oracle@sha256:<digest>" # placeholder
  licenses:
    - id: "etherscan_terms_v2025-01"
    - id: "beacon_api_terms_v2025-02"

```

---

## References & Availability

- Triadic Self-Coherence (TSC) framework (2024). Mathematical foundations and axioms.
- Coherent Blockchain (2025). Protocol direction and leverage-based gating.
- Incident reports: Terra/Luna post-mortems, The DAO reentrancy analyses, archival Bitcoin studies.
- Earth Mover's Distance: Rubner, Y., Tomasi, C., & Guibas, L. J. (2000). The Earth Mover's Distance as a metric for image retrieval. *International Journal of Computer Vision*.
- Edit distance: Wagner, R. A., & Fischer, M. J. (1974). The string-to-string correction problem. *Journal of the ACM*.

A full bibliography with DOIs/URLs will be included in the arXiv submission.

---

## ArXiv Submission Checklist

- **Categories:** cs.CR (Cryptography and Security), cs.DC (Distributed Computing)
- **Keywords:** blockchain measurement, coherence, DeFi risk, reproducibility, zero-knowledge
- **Source:** LaTeX or PDF from Markdown
- **Include:** .bib with 5-8 core references (TSC framework, Terra/DAO post-mortems, EMD/edit-distance citations)
- **Artifacts:** Link to validation notebooks repo when ready
- **Diagrams:** Replace ASCII diagrams with vector PDFs if possible
- **Math rendering:** Uses Unicode symbols for GitHub; convert to LaTeX for arXiv
- **Placeholders:** Clearly marked in schema (e.g., `<eth-rpc>`, `<commit>`, `<blob>`, `<digest>`)

---

## Conclusion

Phase-1 specification provides a complete framework for measuring blockchain coherence. The validation notebooks (Month 1-2) will prove whether this framework works as designed. If validation succeeds, we proceed to build the oracle infrastructure. If validation reveals issues, we iterate on the methodology with clear data about what needs adjustment.

**This honest, empirical approach is how engineering works:** specify, validate, iterate, deploy.

**Next artifact to ship:** Three frozen-input notebooks proving TSC framework can retroactively detect coherence collapse before financial loss:

- `terra_2022-04-15.ipynb` – target  $C_{\Sigma} \approx 0.27 \pm 0.05$
- `dao_2016-06-10.ipynb` – target  $C_{\Sigma} \approx 0.56 \pm 0.05$
- `mtgox_2014-02-10.ipynb` – target  $C_{\Sigma} \approx 0.31 \pm 0.06$

Each outputs:

- `attestation.json` (scores, CI, gates, cause map, B, η)
- `provenance.json` (input hashes, method versions, seeds, container digest, licenses)

---

**End of v1.0.0**